

Python Project 5 - K-Means Clustering

In this project we want to explore a dataset which contains information on breakfast cereals. We are only going to consider the manufacturers, cereal ratings, sugar and carbohydrate levels. Some of the values in sugars and carbohydrate features have invalid data which we must address.

We want to answer the following questions.

1. What cereal manufacturers have the most/least varieties?
2. What cereal has the highest rating?
3. Investigate sugar levels in cereals
 - Does each manufacturer provide cereals equally among the sugar levels?
 - How are the cereals clustered by sugar levels (low, medium, high)
 - What are the cereals which have the highest sugar level? lowest sugar levels?
 - How does your favorite breakfast cereal rank?
1. Investigate carbohydrate levels
 - Cluster cereals in low, medium and high level of carbs
 - What cereals are low sugar and low carbs?
 - What cereals are high sugar and high carbs?

Steps:

- Input libraries and K-Means model from scikit-learn
- Create dataframe and print out a few lines
- Drop all features in dataframe except those we are using: Name, Manufacturer, Sugars and Carbohydrates, and Rating; alternately you can create a new dataframe with just the features you want
- Check to see if there are any NaN values in sugars or carbohydrates features
- Check to see if there are any invalid values (that is values <0) in sugars or carbohydrates features; if there are 2 or less cereals with these invalid values drop the cereal(s) from dataframe; otherwise replace negative values with 0.
- Plot number of cereals for each manufacturer; if any manufacturer has only 1 cereal, drop that data instance. Which manufacturer has the most varieties of cereals.
- In dataframe the manufacturer is listed by a single letter; for plotting purposes add a column giving the actual name and delete feature 'mfr'
- Use a plot to determine which manufacturer has cereals with the highest ratings.
- What cereal has the highest rating?
- Plot sugars levels vs manufacturer and determine which brand has lowest sugar levels.
- Cluster data using K-Means and the sugars feature with clusters low, middle and high sugar levels; print out cluster centroids; add cluster as feature in dataframe.
- Determine which cluster is associated with low, mid or high sugar levels; add feature giving sugar level corresponding to cluster
- Create a plot to show how cereals are distributed among sugar levels; what cereals have highest sugar levels? lowest sugar levels?
- Repeat the steps you did for investigating sugar levels using carbohydrate levels.
- Print out the cereals that are low sugars and low carbohydrates.
- Print out the cereals that are low sugars and high carbohydrates.

Data Set Description

This dataset contains information on 77 different breakfast cereals. The features are

1. Cereal Name

2. Manufacturer
- A → American Home Food Products

• G → General Mills

• K → Kelloggs

• N → Nabisco

• P → Post

• Q → Quaker Oats

• R → Ralston Purina
1. Type (C → Cold, H → Hot)

2. Calories (per serving)

3. Protein (in grams)

4. Fat (in grams)

5. Sodium (in milligrams)

6. Fiber (in grams of dietary fiber)

7. Carbo (grams of complex carbohydrates)

8. Sugars (in grams)

9. Potass (milligrams of potassium)

10. Vitamins (possible values are 0, 25 or 100 indicating percent of FDA recommended vitamins and minerals)

11. Shelf (display shelf with possible values 1, 2, or 3 counting from floor)

12. Weight (weight of one serving in ounces)




13. Cups (number of cups in one serving)

14. Rating (Consumer ratings of cereal from 0 to 100)

```
# import libraries and K-Means function

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.cluster import KMeans
```

```
# Create dataframe for data
df = pd.read_csv('cereal.csv')
df = df[['name', 'mfr', 'sugars', 'carbo', 'rating']]
df
```

	name object	mfr object	sugars int64	carbo float64	rating float64	
	100% Bran 1.3% 100% Natura... _ 1.3% 75 others 97.4%	K 29.9% G 28.6% 5 others 41.6%	-1 - 15 	-1.0 - 23.0 	18.042851 - 93.704... 	
0	100% Bran	N	6	5.0	68.402973	
1	100% Natural Bran	Q	8	8.0	33.983679	
2	All-Bran	K	5	7.0	59.425505	
3		K	0	8.0	93.704912	

	All-Bran with Extra					
4	Almond Delight	R	8	14.0	34.384843	
5	Apple Cinnamon Cheerios	G	10	10.5	29.509541	
6	Apple Jacks	K	14	11.0	33.174094	
7	Basic 4	G	8	18.0	37.038562	
8	Bran Chex	R	6	15.0	49.120253	
9	Bran Flakes	P	5	13.0	53.313813	

This chart is empty

```
# drop features we aren't using or create a new dataframe with only the features we want
```

```
# Check to see if any instances have NaN for entries using .info
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0   name    77 non-null        object
1   mfr     77 non-null        object
2   sugars  77 non-null        int64
3   carbo   77 non-null        float64
4   rating  77 non-null        float64
dtypes: float64(2), int64(1), object(2)
memory usage: 3.1+ KB
None
```

```
# Check to see if any cereals have negative values for sugars or carbohydrates; if 2 or less cereals, drop those
# data instances; otherwise replace negative values with 0
```

```
neg_sugars = df[df['sugars'] < 0]
neg_carbs = df[df['carbo'] < 0]
```

```
if len(neg_sugars) <= 2:
    df = df[df['sugars'] >= 0]
if len(neg_carbs) <= 2:
    df = df[df['carbo'] >= 0]
```


```
# Address negative values
```

```
df.loc[df['sugars'] < 0, 'sugars'] = 0
df.loc[df['carbo'] < 0, 'carbo'] = 0
```

```
# set background grid for plots
sns.set_style('whitegrid')
```

```
# Plot number of products per manufacturer
sns.countplot(x='mfr', data=df)
```

```
<AxesSubplot:xlabel='mfr', ylabel='count'>
```



```
# There is only one cereal from American Home Foods Company so we drop that data sample
#
counts = df.mfr.value_counts()

# Type df[df.mfr == 'A'] and it will give you all data from this manufacturer.
# Make sure you set df = drop command; use command to make sure it was deleted
#

drop_mfrs = counts[counts == 1].index
df = df[~df.mfr.isin(drop_mfrs)]

df[df.mfr == 'A']
```

	name object	mfr object	sugars int64	carbo float64	rating float64	

```
# For plots we would like the name of manufacturer instead of just "N" or "Q"
# Use .map or .apply
#

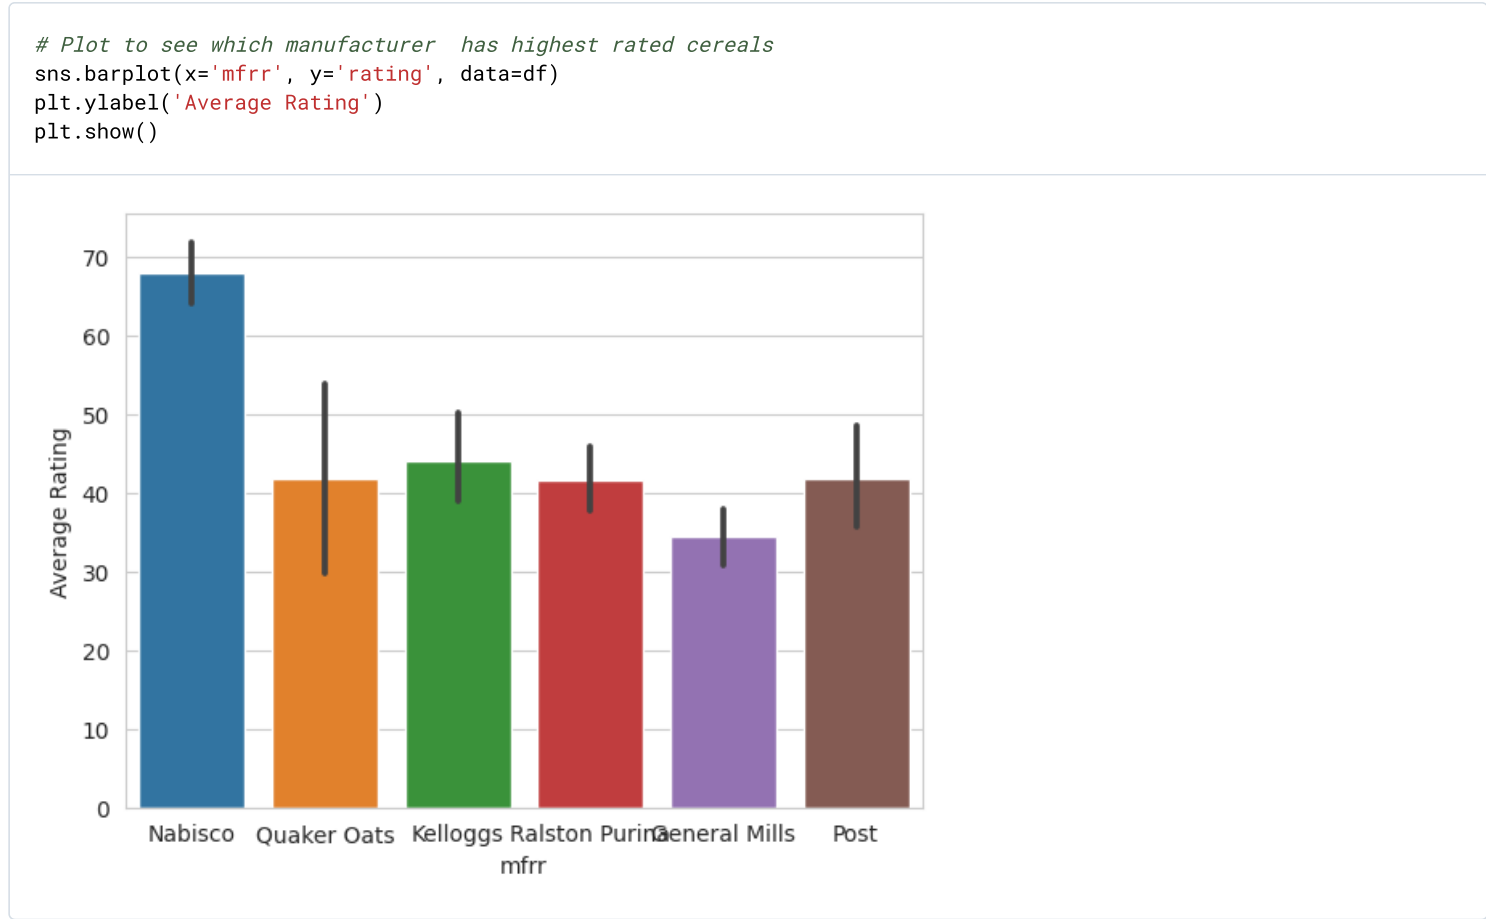
mfr_names = {'A': 'American Home Food Products', 'G': 'General Mills', 'K': 'Kelloggs',
             'N': 'Nabisco', 'P': 'Post', 'Q': 'Quaker Oats', 'R': 'Ralston Purina'}

df['mfrn'] = df['mfr'].map(mfr_names)

# Drop the 'mfr' column
```

df = df.drop('mfr', axis=1)
df

	name object	sugars int64	carbo float64	rating float64	mfr object	
	100% Bran 1.3% 100% Natural Bran 1.3% 73 others 97.3%	0 - 15 	5.0 - 23.0 	18.042851 - 93.704... 	Kelloggs 30.7% General Mills 29.3% 4 others 40%	
0	100% Bran	6	5.0	68.402973	Nabisco	
1	100% Natural Bran	8	8.0	33.983679	Quaker Oats	
2	All-Bran	5	7.0	59.425505	Kelloggs	
3	All-Bran with Extra Fiber	0	8.0	93.704912	Kelloggs	
4	Almond Delight	8	14.0	34.384843	Ralston Purina	
5	Apple Cinnamon Cheerios	10	10.5	29.509541	General Mills	
6	Apple Jacks	14	11.0	33.174094	Kelloggs	
7	Basic 4	8	18.0	37.038562	General Mills	
8	Bran Chex	6	15.0	49.120253	Ralston Purina	
9	Bran Flakes	5	13.0	53.313813	Post	



```
# Which cereal is rated highest?  
print('Nabisco is rated the highest, above Kelloggs and Post')
```

Nabisco is rated the highest, above Kelloggs and Post

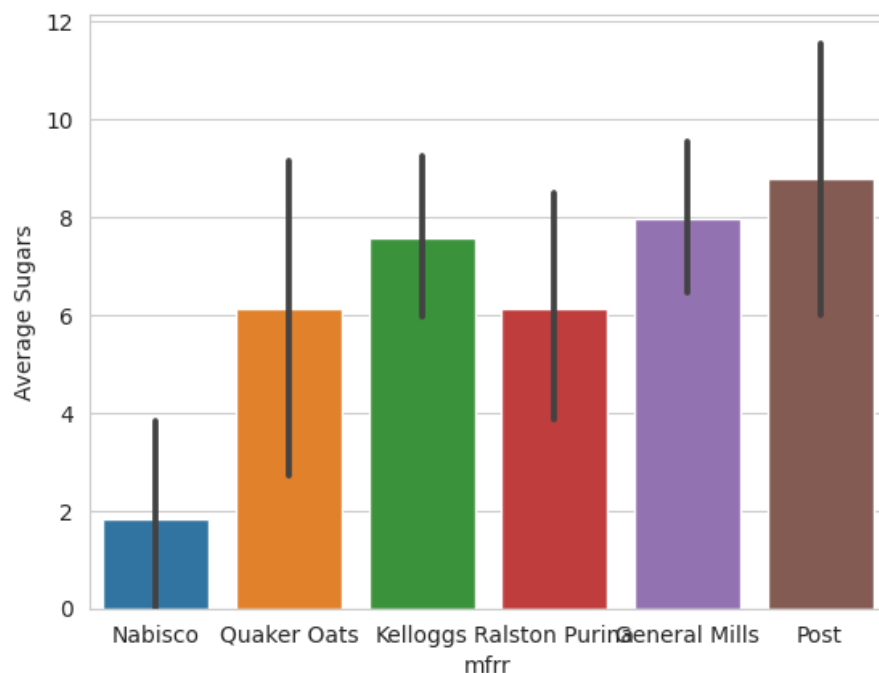
```
# Look at sugars per brand by plotting
sns.barplot(x='mfr', y='sugars', data=df)

plt.ylabel('Average Sugars')
```



Practice Files / Python Project 5 - K-Means Clustering - Duplicate Published at Apr 9, 2023 Unlisted

```
plt.show()
```



Cluster by sugars into highest, middle and lowest levels; random initial guess

```
# Get data for clustering
sugars_df = df[['sugars']]
```

```
# Form model, fit data and print out cluster centers
kmeans = KMeans(n_clusters=3)
kmeans.fit(sugars_df)
print(kmeans.cluster_centers_)
```


```
[[ 1.95833333]
 [12.03846154]
 [ 6.84      ]]
```

```
# Add column to dataframe for this clusters, say sugars_clusters
sugars_clusters = kmeans.predict(sugars_df)
```

```
df['sugars_clusters'] = sugars_clusters
```

```
# Plot clusters
plt.scatter(sugars_df, df['rating'], c=df['sugars_clusters'])
```

<matplotlib.collections.PathCollection at 0x7f91860d1790>



```
# Determine which cluster number corresponds to lowest, middle and highest level and create a new
# column in dataframe using .map
```


```
low_sugar_cluster = 0
middle_sugar_cluster = 1
high_sugar_cluster = 2
```

```
df['sugar_level'] = df['sugars_clusters'].map({low_sugar_cluster: 'low',
                                              middle_sugar_cluster: 'middle',
                                              high_sugar_cluster: 'high'})
```

```
#
```

```
# How are cereals distributed among the 3 levels?
sns.countplot(x='sugar_level', hue='sugars_clusters', data=df)
plt.xlabel('Sugars level')
plt.ylabel('Count of cereals')
plt.title('Distribution of cereals among sugar level clusters')
plt.show()
```

```
# the most cereals sold on average are the ones sold in the middle sugar levels
```



```
# Which cereals have the highest sugar levels
#
```

```
highest_sugar_cereals = df.sort_values('sugars', ascending=False).head(10)
```

```
print(highest_sugar_cereals[['name', 'sugars']])
```

#printed is the top 10 cereals with the highest sugar levels with Smacks, Golden Crisps, and Post Nat. Raisin Bran.

	name	sugars
66	Smacks	15
30	Golden Crisp	15
52	Post Nat. Raisin Bran	14
70	Total Raisin Bran	14
6	Apple Jacks	14
24	Froot Loops	13
46	Mueslix Crispy Blend	13
14	Cocoa Puffs	13
18	Count Chocula	13
10	Cap'n'Crunch	12

```
# Which cereals have the lowest sugar levels
#
```

```
lowest_sugar_cereals = df.sort_values('sugars', ascending=True).head(10)
```

```
print(lowest_sugar_cereals[['name', 'sugars']])
```

#printed are the 10 cereals with the lowest sugar levels with Cream of wheat, All-Bran with extra fiber, and Puffed


	name	sugars
20	Cream of Wheat (Quick)	0
3	All-Bran with Extra Fiber	0
54	Puffed Rice	0
55	Puffed Wheat	0
65	Shredded Wheat spoon size	0
63	Shredded Wheat	0
64	Shredded Wheat 'n'Bran	0
11	Cheerios	1
50	Nutri-grain Wheat	2
61	Rice Chex	2

```
# If you eat a particular cereal like Apple Jacks, Froot Loops, etc. what cluster is it in?
my_cereal = 'Cheerios'
print(f"The data instance and sugar cluster for {my_cereal} is {df[df.name == my_cereal].sugars_clusters}")
```

The data instance and sugar cluster for Cheerios is 11 0
 Name: sugars_clusters, dtype: int32

Repeat calculations and plots for carbohydrates instead of sugars

```
sns.barplot(x='mfrr', y='carbo', data=df)
plt.ylabel('Average Carbohydrates')
plt.show()
```




```
carbo_df = df[['carbo']]
kmeans = KMeans(n_clusters=3)
kmeans.fit(carbo_df)
```

```
print(kmeans.cluster_centers_)
carbo_clusters = kmeans.predict(carbo_df)

df['carbo_clusters'] = carbo_clusters
plt.scatter(carbo_df, df['rating'], c=df['carbo_clusters'])
```

```
[[14.81944444]
 [20.41176471]
 [10.38636364]]
```

<matplotlib.collections.PathCollection at 0x7f91803c5210>



```
low_carbo_cluster = 0
middle_carbo_cluster = 1
high_carbo_cluster = 2

df['carbo_level'] = df['carbo_clusters'].map({low_carbo_cluster: 'low',
                                              middle_carbo_cluster: 'middle',
                                              high_carbo_cluster: 'high'})

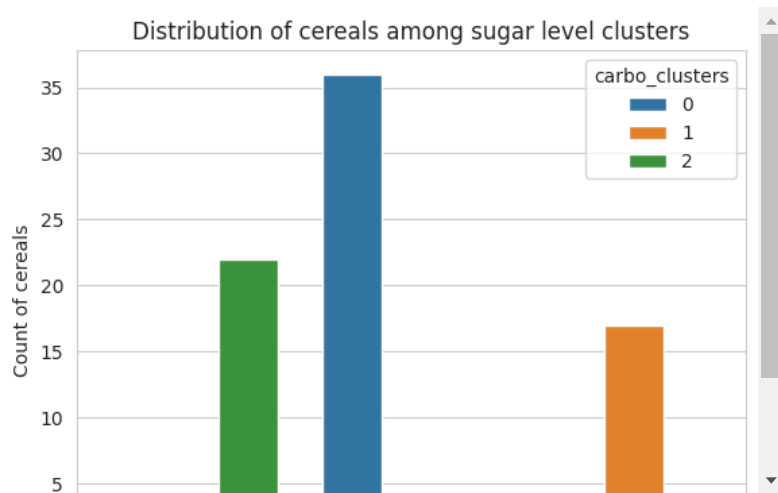
# How are cereals distributed among the 3 levels?
sns.countplot(x='carbo_level', hue='carbo_clusters', data=df)
plt.xlabel('Carbohydrate level')
plt.ylabel('Count of cereals')
plt.title('Distribution of cereals among sugar level clusters')
plt.show()
#Low carbohydrate levels have the highest sales on average and middle with the lowest sales on average

lowest_sugar_cereals = df.sort_values('carbo', ascending=True).head(10)

print(lowest_sugar_cereals[['name', 'carbo']])

highest_sugar_cereals = df.sort_values('carbo', ascending=False).head(10)

print(highest_sugar_cereals[['name', 'carbo']])
```



	name	carbo
0	100% Bran	5.0
2	All-Bran	7.0
1	100% Natural Bran	8.0
3	All-Bran with Extra Fiber	8.0
66	Smacks	9.0
19	Cracklin' Oat Bran	10.0
55	Puffed Wheat	10.0
5	Apple Cinnamon Cheerios	10.5
59	Raisin Nut Bran	10.5
30	Golden Crisp	11.0
61	Rice Chex	23.0
15	Corn Chex	22.0
62	Rice Krispies	22.0
72	Triples	21.0
49	Nutri-Grain Almond-Raisin	21.0
21	Crispix	21.0
40	Kix	21.0
16	Corn Flakes	21.0
20	Cream of Wheat (Quick)	21.0
69	Total Corn Flakes	21.0

What cereals are high carbs and low sugar? What are low carbs and low sugar?

```
high_carbs_low_sugar = df[(df['carbo'] > df['carbo'].mean()) & (df['sugars'] < df['sugars'].mean())]
print(high_carbs_low_sugar[['name', 'carbo', 'sugars']])
```

	name	carbo	sugars
8	Bran Chex	15.0	6
11	Cheerios	17.0	1
15	Corn Chex	22.0	3
16	Corn Flakes	21.0	2

20	Cream of Wheat (Quick)	21.0	0
21	Crispix	21.0	3
23	Double Chex	18.0	5
32	Grape Nuts Flakes	15.0	5
33	Grape-Nuts	17.0	3
38	Just Right Crunchy Nuggets	17.0	6
40	Kix	21.0	3
47	Multi-Grain Cheerios	15.0	6
49	Nutri-Grain Almond-Raisin	21.0	7
50	Nutri-grain Wheat	18.0	2
53	Product 19	20.0	3
60	Raisin Squares	15.0	6
61	Rice Chex	23.0	2
62	Rice Krispies	22.0	3
63	Shredded Wheat	16.0	0
64	Shredded Wheat 'n'Bran	19.0	0
65	Shredded Wheat spoon size	20.0	0
67	Special K	16.0	3
68	Strawberry Fruit Wheats	15.0	5
69	Total Corn Flakes	21.0	3
71	Total Whole Grain	16.0	3
72	Triples	21.0	3
74	Wheat Chex	17.0	3
75	Wheaties	17.0	3

```
low_carbs_low_sugar = df[(df['carbo'] < df['carbo'].mean()) & (df['sugars'] < df['sugars'].mean())]
print(low_carbs_low_sugar[['name', 'carbo', 'sugars']])
```

	name	carbo	sugars
0	100% Bran	5.0	6
2	All-Bran	7.0	5
3	All-Bran with Extra Fiber	8.0	0
9	Bran Flakes	13.0	5
13	Clusters	13.0	7
19	Cracklin' Oat Bran	10.0	7
26	Frosted Mini-Wheats	14.0	7
34	Great Grains Pecan	13.0	4
41	Life	12.0	6
54	Puffed Rice	13.0	0
55	Puffed Wheat	10.0	0
56	Quaker Oat Squares	14.0	6